

```
In [18]: #Load Libraries
import os
import numpy as np
import torch
import glob
import torch.nn as nn
from torchvision.transforms import transforms
from torch.utils.data import DataLoader
from torch.optim import Adam
from torch.autograd import variable
import torchvision
import pathlib
```

```
In [19]: #checking for device
device=torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
In [3]: print(device)
```

cpu

```
In [4]: #Transforms
transformer=transforms.Compose([
    transforms.Resize((150,150)),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(), #0-255 to 0-1, numpy to tensors
    transforms.Normalize([0.5,0.5,0.5], # 0-1 to [-1,1] , formula (x-mean)/std
                          [0.5,0.5,0.5])
])
```

```
In [5]: #DataLoader

#Path for training and testing directory
train_path='E:\CNN for Bone fracture Image 02 -\Fracture detection\seg_train\seg_train'
test_path='E:\CNN for Bone fracture Image 02 -\Fracture detection\seg_test\seg_test'

train_loader=DataLoader(
    torchvision.datasets.ImageFolder(train_path,transform=transformer),
    batch_size=256, shuffle=True
)
test_loader=DataLoader(
    torchvision.datasets.ImageFolder(test_path,transform=transformer),
    batch_size=128, shuffle=True
)
```

```
In [6]: #categories
root=pathlib.Path(train_path)
classes=sorted([j.name.split('/')[-1] for j in root.iterdir()])
```

```
In [7]: print(classes)

['Fracture Image', 'Non-fracture Image']
```

```
In [8]: class ConvNet(nn.Module):
    def __init__(self,num_classes=6):
        super(ConvNet,self).__init__()
```

```

#Output size after convolution filter
#((w-f+2P)/s) +1

#Input shape= (256,3,150,150)

self.conv1=nn.Conv2d(in_channels=3,out_channels=12,kernel_size=3,stride=1,pad
#Shape= (256,12,150,150)
self.bn1=nn.BatchNorm2d(num_features=12)
#Shape= (256,12,150,150)
self.relu1=nn.ReLU()
#Shape= (256,12,150,150)

self.pool=nn.MaxPool2d(kernel_size=2)
#Reduce the image size be factor 2
#Shape= (256,12,75,75)

self.conv2=nn.Conv2d(in_channels=12,out_channels=20,kernel_size=3,stride=1,pad
#Shape= (256,20,75,75)
self.bn2=nn.BatchNorm2d(num_features=20)
#Shape= (256,20,75,75)
self.relu2=nn.ReLU()
#Shape= (256,20,75,75)

self.pool2=nn.MaxPool2d(kernel_size=2)
#Reduce the image size be factor 2
#Shape= (256,20,75,75)

self.conv3=nn.Conv2d(in_channels=20,out_channels=32,kernel_size=3,stride=1,pad
#Shape= (256,32,75,75)
self.bn3=nn.BatchNorm2d(num_features=32)
#Shape= (256,32,75,75)
self.relu3=nn.ReLU()
#Shape= (256,32,75,75)

self.pool3=nn.MaxPool2d(kernel_size=2)
#Reduce the image size be factor 2
#Shape= (256,32,75,75)

self.fc=nn.Linear(in_features=75 * 75 * 32,out_features=num_classes)

#Feed forward function

def forward(self,input):
    output=self.conv1(input)
    output=self.bn1(output)
    output=self.relu1(output)

    output=self.pool(output)

    output=self.conv2(output)
    output=self.bn2(output)
    output=self.relu2(output)

```

```

output=self.conv3(output)
output=self.bn3(output)
output=self.relu3(output)

```

*#Above output will be in matrix form, with shape (256,32,75,75)*

```
output=output.view(-1,32*75*75)
```

```
output=self.fc(output)
```

```
return output
```

```
In [9]: model=ConvNet(num_classes=2).to(device)
```

```
In [10]: #Optimizer and loss function
optimizer=Adam(model.parameters(),lr=0.001,weight_decay=0.0001)
loss_function=nn.CrossEntropyLoss()
```

```
In [11]: num_epochs=8
```

```
In [12]: #calculating the size of training and testing images
train_count=len(glob.glob(train_path+'**/*.jpg'))
test_count=len(glob.glob(test_path+'**/*.jpg'))
```

```
In [13]: print(train_count,test_count)
```

```
9193 8907
```

```
In [14]: #Model training and saving best model

best_accuracy=0.0

for epoch in range(num_epochs):

    #Evaluation and training on training dataset
    model.train()
    train_accuracy=0.0
    train_loss=0.0

    for i, (images,labels) in enumerate(train_loader):
        if torch.cuda.is_available():
            images=Variable(images.cuda())
            labels=Variable(labels.cuda())

        optimizer.zero_grad()

        outputs=model(images)
        loss=loss_function(outputs,labels)
        loss.backward()
        optimizer.step()

        train_loss+= loss.cpu().data*images.size(0)
        _,prediction=torch.max(outputs.data,1)
```

```

        train_accuracy+=int(torch.sum(prediction==labels.data))

train_accuracy=train_accuracy/train_count
train_loss=train_loss/train_count

# Evaluation on testing dataset
model.eval()

test_accuracy=0.0
for i, (images,labels) in enumerate(test_loader):
    if torch.cuda.is_available():
        images=Variable(images.cuda())
        labels=Variable(labels.cuda())

    outputs=model(images)
    _,prediction=torch.max(outputs.data,1)
    test_accuracy+=int(torch.sum(prediction==labels.data))

test_accuracy=test_accuracy/test_count

print('Epoch: '+str(epoch)+' Train Loss: '+str(train_loss)+' Train Accuracy: '+str

#Save the best model
if test_accuracy>best_accuracy:
    torch.save(model.state_dict(),'best_checkpoint.model')
    best_accuracy=test_accuracy

```

```

Epoch: 0 Train Loss: tensor(8.1191) Train Accuracy: 0.5823996519090613 Test Accuracy:
0.6608285618053217
Epoch: 1 Train Loss: tensor(0.8034) Train Accuracy: 0.7579680191450017 Test Accuracy:
0.8469742898843606
Epoch: 2 Train Loss: tensor(0.3120) Train Accuracy: 0.8742521483737626 Test Accuracy:
0.7675985180195352
Epoch: 3 Train Loss: tensor(0.2583) Train Accuracy: 0.8952463831175894 Test Accuracy:
0.9500392949365667
Epoch: 4 Train Loss: tensor(0.2048) Train Accuracy: 0.9233112150549331 Test Accuracy:
0.9575614685079151
Epoch: 5 Train Loss: tensor(0.1680) Train Accuracy: 0.9366909605134341 Test Accuracy:
0.938026271471876
Epoch: 6 Train Loss: tensor(0.1052) Train Accuracy: 0.9654084629609485 Test Accuracy:
0.9811384304479622
Epoch: 7 Train Loss: tensor(0.0845) Train Accuracy: 0.9749809637767867 Test Accuracy:
0.9850679241046368

```